

Large-Scale and Multi-Structured Databases

Project Design

BeatBuddy



Luca Arduini
Giovanni Enrico Loni
Lorenzo Mancinelli

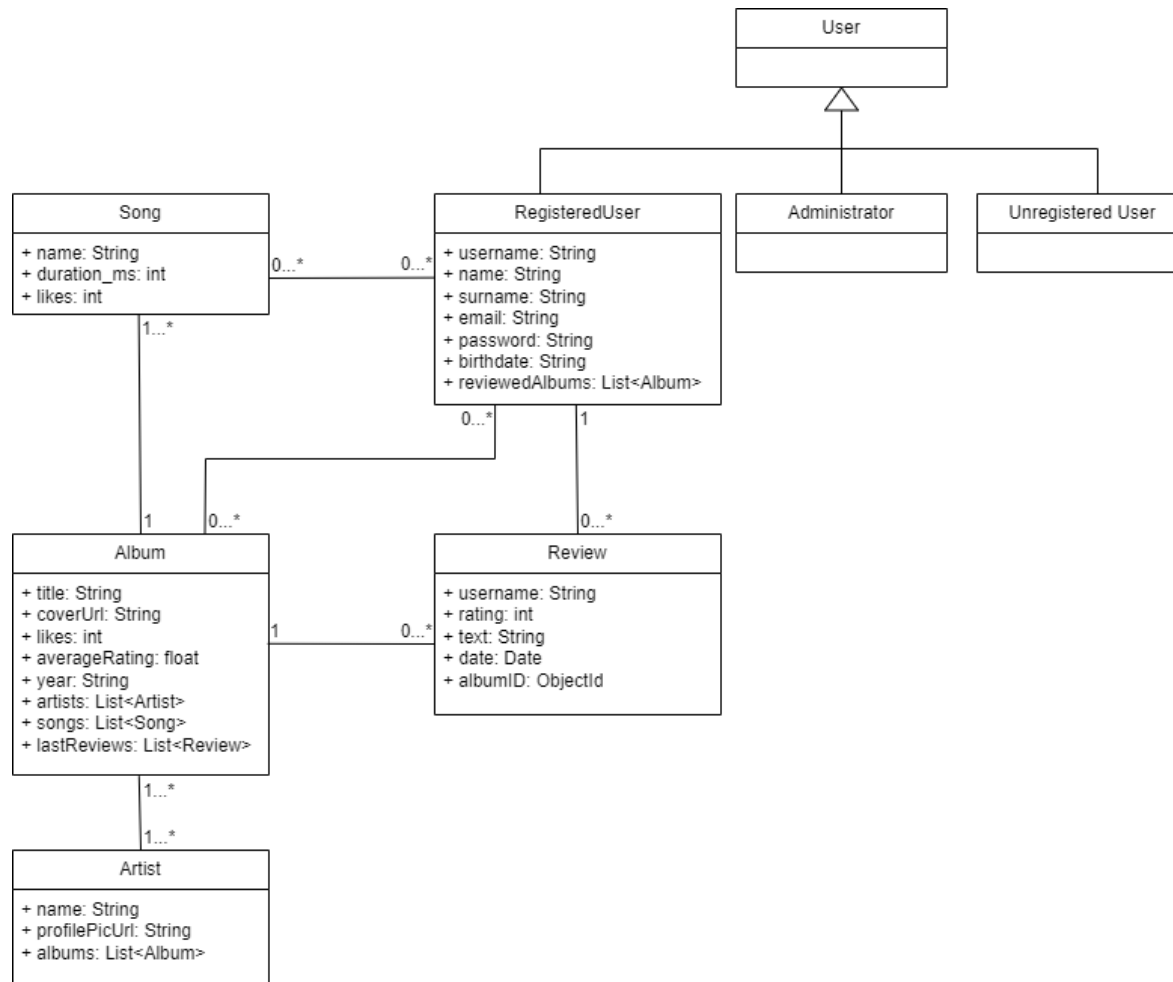
Application Highlights

BeatBuddy merges the world of social networking with a passion for music, highlighting two key features:

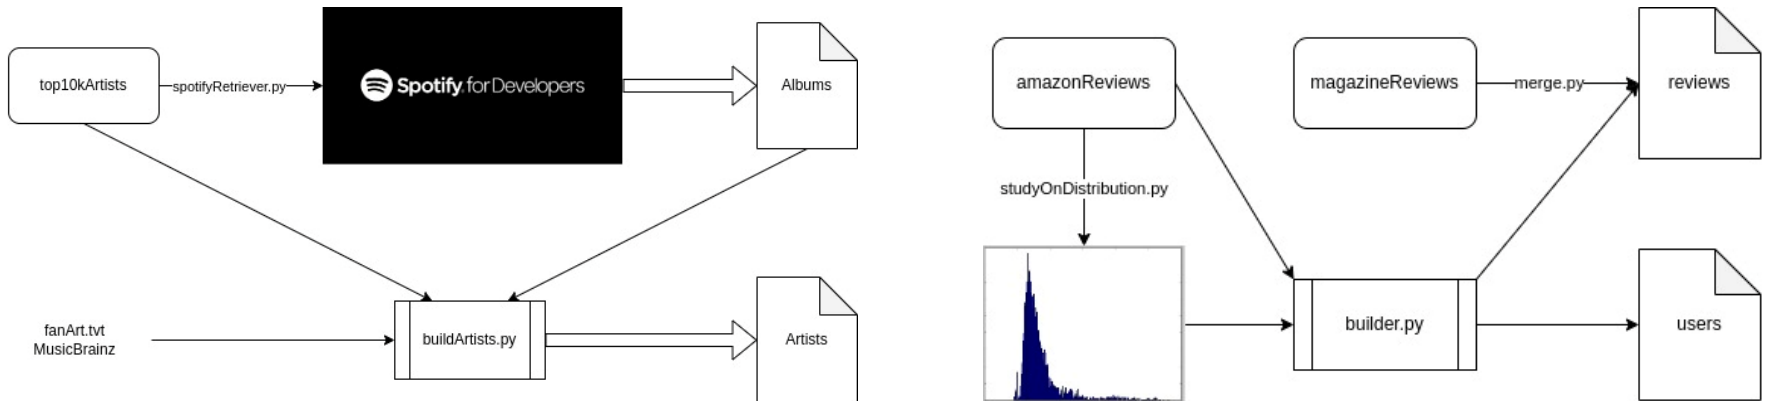
- **Search:** Our users can effortlessly explore albums, songs, and artists. Each search leads to a dedicated page brimming with details such as track listings, related album collections, and reviews from other music lovers.
- **Discover:** We're all about helping you find fresh sounds. The app shines a spotlight on trending tracks and offers personalized recommendations, either tailored to your own taste or influenced by your friends' musical preferences.

And behind the scenes, our admins keep the rhythm going with powerful analytics, staying in tune with how the app is used and ensuring the best experience for everyone.

UML Class Diagram



Dataset Description



albums

Storage size:	Documents:	Avg. document size:	Indexes:	Total index size:
94.78 MB	53 K	2.71 kB	3	23.88 MB

artists

Storage size:	Documents:	Avg. document size:	Indexes:	Total index size:
3.59 MB	7.9 K	960.00 B	1	442.37 kB

reviews

Storage size:	Documents:	Avg. document size:	Indexes:	Total index size:
89.37 MB	160 K	744.00 B	1	4.30 MB

users

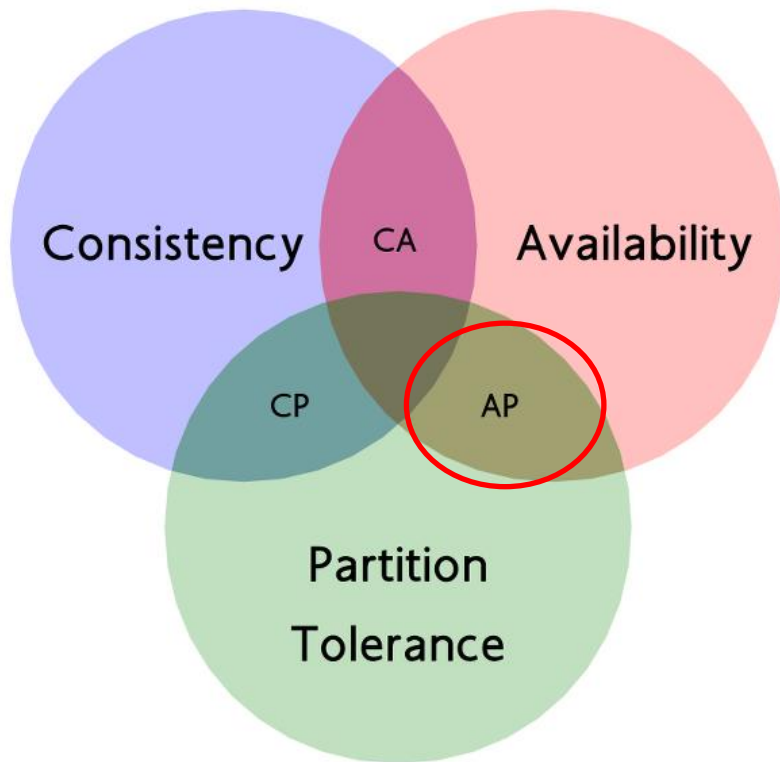
Storage size:	Documents:	Avg. document size:	Indexes:	Total index size:
21.67 MB	57 K	674.00 B	2	3.26 MB

```
1.5M data/databases/beatbuddy/schema/index/range-1.0/8
1.9M data/databases/beatbuddy/schema/index/range-1.0/9
25M data/databases/beatbuddy/schema/index/range-1.0/10
29M data/databases/beatbuddy/schema/index/range-1.0
14M data/databases/beatbuddy/schema/index/token-lookup-1.0/2
1.5M data/databases/beatbuddy/schema/index/token-lookup-1.0/1
15M data/databases/beatbuddy/schema/index/token-lookup-1.0
43M data/databases/beatbuddy/schema/index
43M data/databases/beatbuddy/schema
764M data/databases/beatbuddy/
764M total
```

Non-Functional Requirements

- Implementation as a web application.
- Minimization of single points of failure.
- Emphasis on high availability, even with occasionally outdated data.
- Use of Object-Oriented Programming languages for code development.
- Tolerance for data loss.
- Encryption of user passwords.

CAP Theorem Issue



1. The application needs to be available 24/7
2. We need to avoid a single point of failure
3. We accept potentially out-of-date version of data



WE STAY ON THE AP SIDE OF THE TRIANGLE

MongoDB design

Album's collection

```
{
  "_id": ObjectId('65ab9d61186280a9f3eca093'),
  "artists": Array (1)
    0: "Siriusmo"
  "averageRating": 4.5
  "coverURL": "https://i.scdn.co/image/ab67616d0000b27392df348d788a14c9163d02f2"
  "last_reviews": Array (2)
    0: Object
    1: Object
  "likes": 29
  "songs": Array (7)
    0: Object
      "duration_ms": 195866
      "likes": 5
      "name": "The Plasterer of Love"
    1: Object
    2: Object
    3: Object
    4: Object
    5: Object
    6: Object
  "title": "The Plasterer of Love (Deluxe Edition)"
  "year": "2010"
}
```

User's collection

```
{
  "_id": ObjectId('65ab9d44186280a9f3ebc2bc'),
  "name": "Bianca"
  "surname": "Serlupi"
  "username": "John Hopkins"
  "password": "15a3bfee7081959e1223f95725c77ea5f37475f6e79c3ddd531f67d456e30881"
  "birthDate": "1975-08-06"
  "email": "johnhopkins-bianca@Morpurgo.com"
  "reviewedAlbums": Array (2)
    0: Object
      "artist": "Brian Cross"
      "rating": 5
      "coverUrl": "https://i.scdn.co/image/ab67616d0000b273239c0adf89ebf850d44e4983"
      "albumTitle": "Crossing Lines"
    1: Object
}
```

Review's collection

```
{
  "_id": ObjectId('65ab9d80186280a9f3ed6fc3'),
  "rating": 4
  "text": "The cover art alone of this album is very amusing. It looks like a cer..."
  "albumID": ObjectId('65ab9d71186280a9f3ed0a25')
  "username": "Andre S. Grindle"
  "date": 2024-01-04T21:27:20.000+00:00
}
```

Artist's collection

```
{
  "_id": ObjectId('65afcd0328ee62efd4c52617'),
  "name": "Mudimbi"
  "albums": Array (2)
    0: Object
      "title": "Michel"
      "coverURL": "https://i.scdn.co/image/ab67616d0000b273b2ae826d976a732cd13da211"
    1: Object
  "profilePicUrl": "https://i.pinimg.com/564x/1d/04/a8/1d04a87b8e6cf2c3829c7af2eccf6813.jp..."
}
```

Relevant MongoDB queries

Aggregations:

- **getAlbumsWithMinReviewsByAvgRating_AllTime**

Compiles the ranking of albums with the highest average rating.

- **getSongsByLikes_AllTime**

Creates the list of the most-liked songs of all time

- **getAverageRatingForRecentReviews**

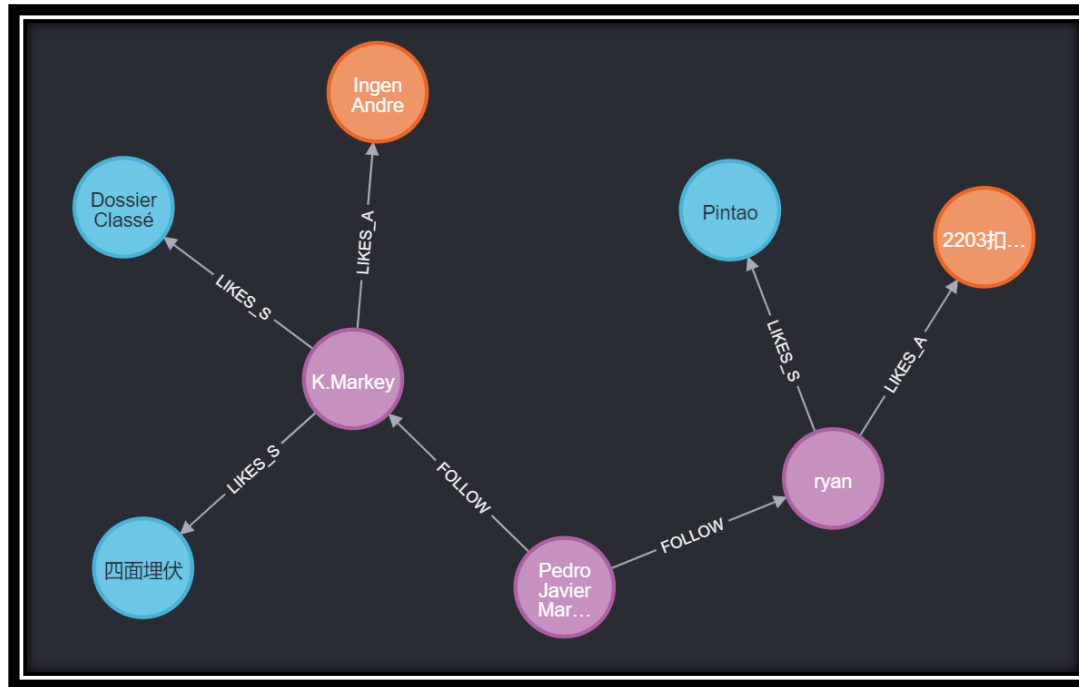
Calculates the average rating for albums that have received a review in the last 24 hours.

```
db.reviews.aggregate([
  {
    $group: { _id: "$albumID", reviewCount: { $sum: 1 } }
  },
  {
    $match: { reviewCount: { $gte: 10 } }
  },
  {
    $lookup: { from: "albums", localField: "_id",
      foreignField: "_id", as: "albumDetails" }
  },
  {
    $project: {
      id: "$_id",
      ...
      coverURL: "$albumDetails.coverURL",
      averageRating: "$albumDetails.averageRating",
    }
  },
  {
    $sort: { averageRating: -1 }
  },
  {
    $limit: 10
  }
]);
```

Relevant operations:

- Display album detail page
- Add a new review
- Search for a song by substring

Neo4J design



Entities:

- **User** (username)
- **Album** (albumName, artistName, coverURL)
- **Song** (albumName, artistName, coverUrl, songName)

Relationships:

- **FOLLOW** (User-User)
- **LIKES_A** (User-Album)
- **LIKES_S** (User-Song)

Relevant Neo4j queries

SUGGESTIONS:

Users:

- Based on the *users already followed*.

Songs:

- Based on *favourite songs among the followed users*.
- **Based on user's liked songs**

```
private ArrayList<Song_Neo4j> findSuggestedSongs_ByTaste(String username) {
    String cypherQuery = "MATCH (targetUser:User {username: $username})-[:LIKES_S]->(likedSong:Song) " +
        "MATCH (similarUser:User)-[:LIKES_S]->(likedSong) " +
        "WHERE targetUser <> similarUser " +
        "MATCH (similarUser)-[:LIKES_S]->(recommendedSong:Song) " +
        "WHERE NOT (targetUser)-[:LIKES_S]->(recommendedSong) " +
        "WITH recommendedSong, COUNT(*) AS recommendationStrength " +
        "RETURN recommendedSong.songName AS songName, " +
        "recommendedSong.albumName AS albumName, " +
        "recommendedSong.artistName AS artistName, " +
        "recommendedSong.coverUrl AS coverUrl " +
        "ORDER BY recommendationStrength " +
        "LIMIT 10";

    return Song_Neo4j.getSongNeo4js(username, cypherQuery, neo4jClient);
}
```

STATISTISC:

- Count all-time likes (Songs and Albums)
- **Count likes in the past week** (Songs and **Albums**)
- Count all-time likes for entities that received likes in the past day

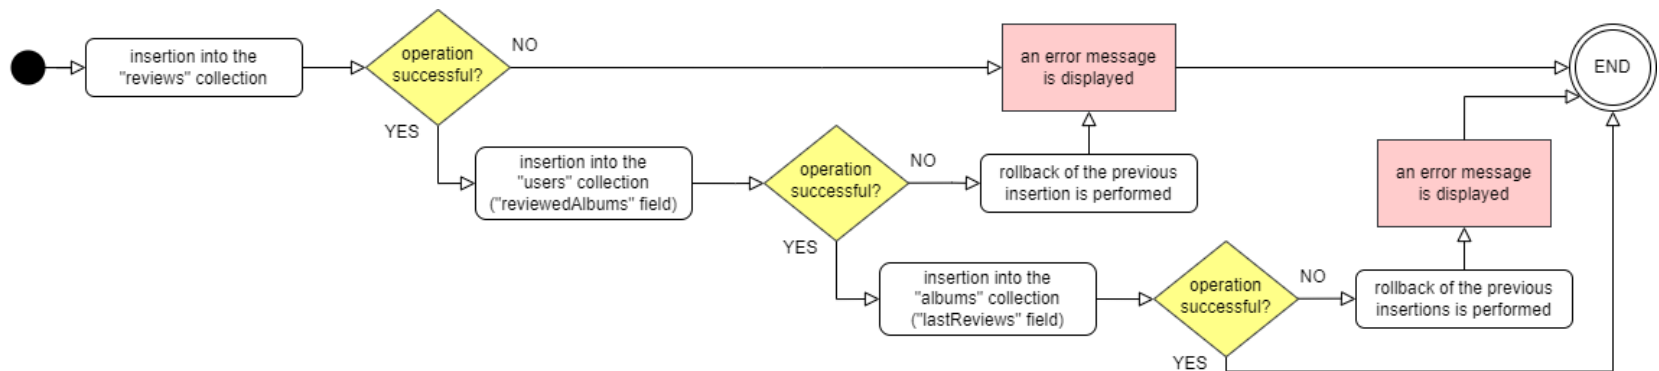
```
public List<AlbumWithLikes> getAlbumsByLikes_LastWeek() {
    String cypherQuery = "MATCH (a:Album) <-[:LIKES_A]- (:User) " +
        "WHERE date(r.timestamp) >= date()- duration('P7D') " +
        "WITH a, count(r) as likes " +
        "RETURN a.albumName AS albumName, a.artistName AS artistName, " +
        "a.coverURL AS coverURL, likes " +
        "ORDER BY likes DESC " +
        "LIMIT 10";

    return AlbumWithLikes.getAlbumWithLikes(cypherQuery, neo4jClient);
}
```

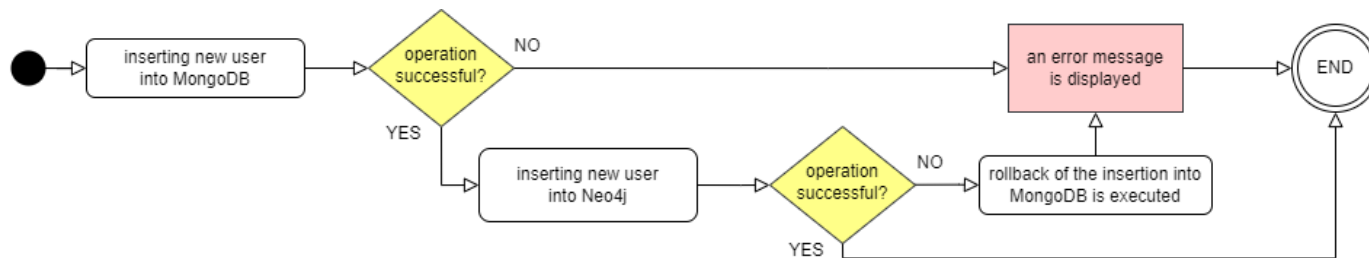
Data Consistency

We use Spring Boot's `@Transactional` to ensure consistency in multiple insertions. If an insertion fails, previous ones in the transaction are automatically rolled back. Let's see how it works in action.

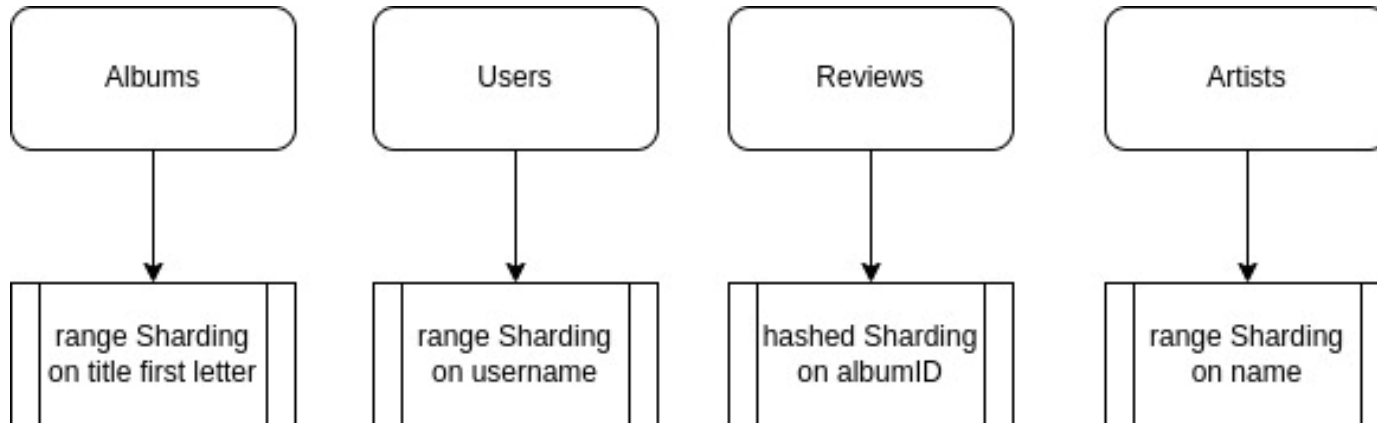
- example 1: Insertion of a new review → three insertions in MongoDB



- example 2: Registration of a new user → one insertion in MongoDB and one in Neo4j

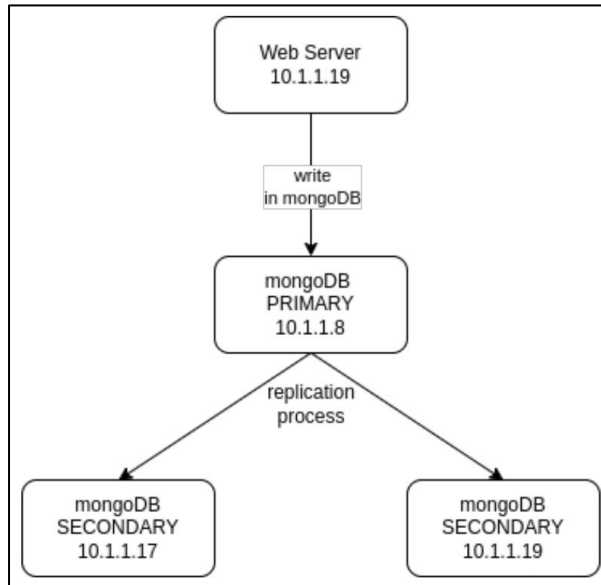


Data Sharding Proposal



Software and Hardware Architecture

MongoDB Write Concern = 1



Neo4J

